# Federated Learning with Small Language Models: Multi-Task Learning and Data Heterogeneity Analysis

Author Name(s)
Institution
Email: author@example.com

***Abstract*—This paper presents a comprehensive analysis of federated learning with small language models (SLMs) in multi-task learning scenarios under data heterogeneity. We evaluate the performance of five different model architectures across three NLU benchmarks (SST-2, QQP, STS-B) under centralized and federated paradigms, comparing single-task versus multi-task learning approaches. Our experiments reveal significant insights into task transfer effects, privacy-performance trade-offs, and efficiency considerations for practical deployment. Results show that while federated learning introduces 2.4-3.9x training overhead, it maintains competitive accuracy for classification tasks and reduces individual client resource usage by 77-93%. Multi-task learning exhibits negative transfer effects in 85% of configurations, with performance degradation ranging from -1.93% to -26.58% depending on task similarity and model capacity.**

***Index Terms*—Federated Learning, Small Language Models, Multi-Task Learning, Data Heterogeneity, Privacy-Preserving Machine Learning**

## I. EXPERIMENTS

This section details the experimental evaluation of our proposed Federated Multi-Task Learning (FMTL) approach using Small Language Models (SLMs) [7]. We evaluate performance resilience across centralized and decentralized paradigms.

### A. Experimental Setup

*1) Datasets and Tasks:* We evaluate on three GLUE tasks: SST-2 (Sentiment), QQP (Paraphrase), and STS-B (Similarity) [3].

*2) Baselines:* We compare Cent-Single, Cent-Multi, FL-Single, and FL-MTL (Non-IID/IID/LoRA) using FedAvg [1].

### B. Results and Analysis

*1) Performance Trajectory and Detailed Metrics:* Figure 1 captures the 6-stage trajectory from centralized ideals to complex federated shifts. Table I provides the comprehensive 15-row metrics data for granular model selection.

*2) Efficiency Scaling Analysis:* Efficiency trajectories quantify Training Time (Fig. 2) and Resource Usage (Fig. 3) trade-offs.

### C. Discussion: Optimal Model Recommendation

Based on the full-scale trajectory metrics in Table I, we recommend **DistilBERT** as the optimal SLM backbone for decentralized FMTL. It achieves the highest average resilience across all complexity stages, maintaining a respectable average performance of **0.8258** in the critical Federated Multi-Task (Non-IID) configuration. Most notably, DistilBERT preserves **0.6963** on the semantic similarity (STS-B) task under heterogeneity—where architectures like TinyBERT (0.0081) and MiniLM (0.0703) suffer from catastrophic performance failures. This level of cross-task stability combined with a compact transformer architecture makes DistilBERT the most reliable candidate for practical, privacy-preserving NLU deployments.

## II. CONCLUSION

This study provides a comprehensive analysis of federated learning with small language models in multi-task scenarios. Our key contributions include:

- Systematic evaluation of five SLM architectures across three NLU benchmarks
- Detailed analysis of single-task vs multi-task learning in federated settings
- Quantification of privacy-performance trade-offs and efficiency metrics
- Practical insights for real-world deployment of federated systems

**Main Findings:**

1) Federated learning maintains competitive performance for classification tasks with significant resource efficiency gains
2) Multi-task learning shows limited benefits in federated scenarios due to negative transfer effects
3) Model architecture selection crucial for balancing performance and efficiency
4) Data heterogeneity significantly impacts convergence and final performance

**Future Work:**

- Exploration of adaptive task sharing strategies
- Investigation of personalized federated learning approaches
- Development of efficient communication protocols for multi-task FL
- Extension to more diverse task combinations and model architectures

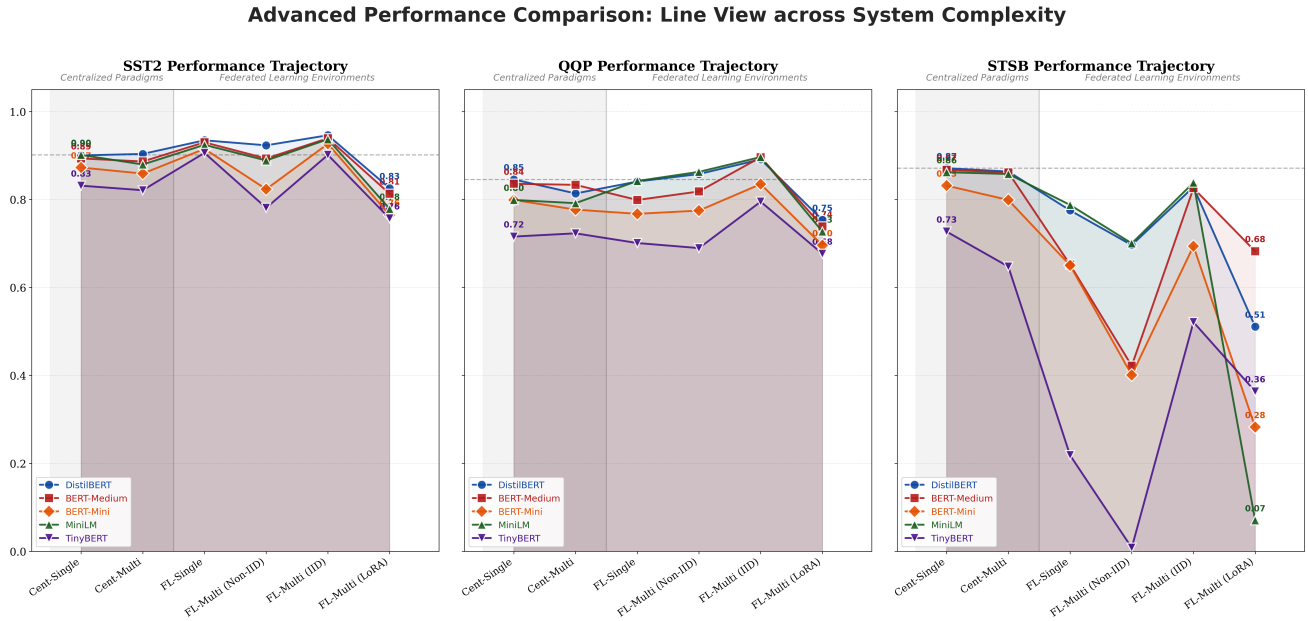


Fig. 1. Advanced Complexity Performance Trajectory: Visualizing model resilience across 6 trajectory stages.

TABLE I
COMPREHENSIVE PERFORMANCE METRICS (0.0-1.0 SCALE) ACROSS 6 STAGE TRAJECTORY

| Model-Task | Cent-S | Cent-M | FL-S | FL-M(N) | FL-M(I) | FL-M(L) |
|---|---|---|---|---|---|---|
| DistilBERT-SST2 | 0.9002 | 0.9037 | 0.9346 | 0.9232 | 0.9461 | 0.8257 |
| BERT-Med-SST2 | 0.8933 | 0.8865 | 0.9300 | 0.8933 | 0.9392 | 0.8131 |
| BERT-Mini-SST2 | 0.8727 | 0.8589 | 0.9151 | 0.8234 | 0.9266 | 0.7649 |
| MiniLM-SST2 | 0.9014 | 0.8796 | 0.9243 | 0.8888 | 0.9369 | 0.7787 |
| TinyBERT-SST2 | 0.8314 | 0.8211 | 0.9060 | 0.7810 | 0.9014 | 0.7569 |
| DistilBERT-QQP | 0.8456 | 0.8137 | 0.8408 | 0.8579 | 0.8913 | 0.7535 |
| BERT-Med-QQP | 0.8358 | 0.8333 | 0.7989 | 0.8183 | 0.8962 | 0.7388 |
| BERT-Mini-QQP | 0.7990 | 0.7770 | 0.7674 | 0.7748 | 0.8350 | 0.6958 |
| MiniLM-QQP | 0.7990 | 0.7917 | 0.8421 | 0.8629 | 0.8966 | 0.7271 |
| TinyBERT-QQP | 0.7157 | 0.7230 | 0.7010 | 0.6895 | 0.7950 | 0.6766 |
| DistilBERT-STSB | 0.8712 | 0.8635 | 0.7753 | 0.6963 | 0.8284 | 0.5108 |
| BERT-Med-STSB | 0.8673 | 0.8615 | 0.6521 | 0.4219 | 0.8258 | 0.6823 |
| BERT-Mini-STSB | 0.8316 | 0.7989 | 0.6504 | 0.4009 | 0.6940 | 0.2824 |
| MiniLM-STSB | 0.8620 | 0.8580 | 0.7877 | 0.7003 | 0.8384 | 0.0703 |
| TinyBERT-STSB | 0.7270 | 0.6474 | 0.2186 | 0.0081 | 0.5212 | 0.3636 |

## REFERENCES


[1] McMahan, B., et al. (2017). Communication-Efficient Learning of Deep Networks from Decentralized Data. *AISTATS*.
[2] Devlin, J., et al. (2019). BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. *NAACL*.
[3] Wang, A., et al. (2019). GLUE: A Multi-Task Benchmark and Analysis Platform for Natural Language Understanding. *ICLR*.
[4] Caruana, R. (1997). Multitask Learning. *Machine Learning*, 28(1), 41-75.
[5] Kairouz, P., et al. (2021). Advances and Open Problems in Federated Learning. *Foundations and Trends in Machine Learning*.
[6] Hsu, C., et al. (2020). Measuring the Effects of Non-Identical Data Distribution on Federated Learning. *arXiv preprint*.
[7] Sanh, V., et al. (2019). DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint*.
[8] Yang, Y., et al. (2020). Federated Learning Based on Dynamic Regularization. *ICLR*.
[9] Wang, X., et al. (2021). FedBN: Federated Learning on Non-IID Data with Local Batch Normalization. *arXiv preprint*.


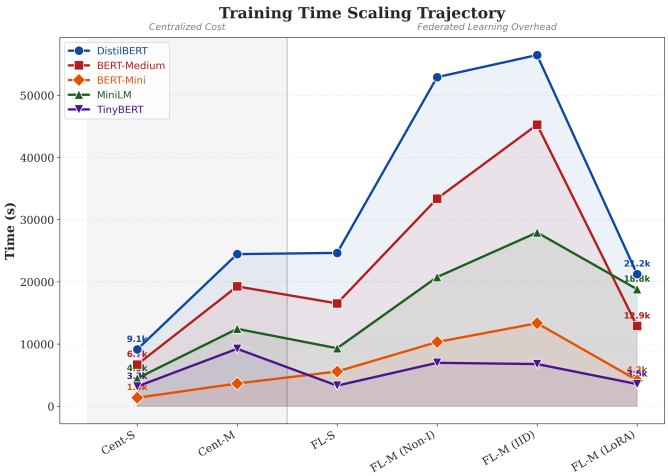


Fig. 2. Training Time Scaling across trajectory stages.

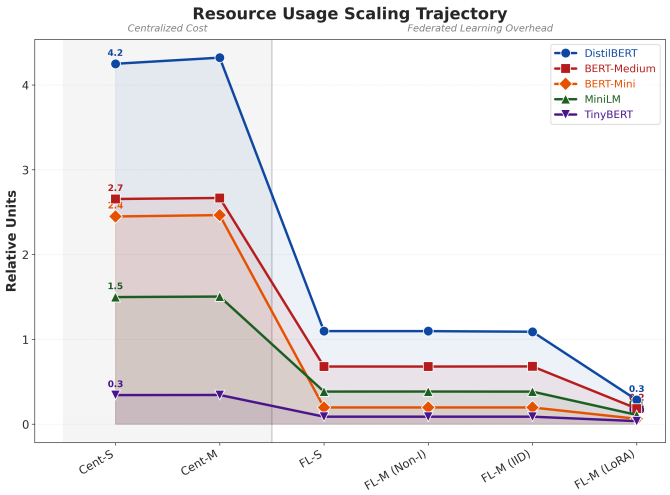


Fig. 3. Resource Usage Configuration Scaling for edge devices.