A Real-Time Federated Multi-Task Learning Framework with Small Language Models for Natural Language Understanding Tasks

Quang-Vinh Pham
Department of Data Science
TMA Tech Group, Gia Lai, Vietnam
Email: pqvinh@tma.com.vn

Quang-Hung Le
Department of Information Technology
Quy Nhon University, Vietnam
Email: lequanghung@qnu.edu.vn

***Abstract*—The adoption of Federated Learning (FL) for privacy-preserving Natural Language Understanding (NLU) is often hindered by the resource constraints of edge devices and the complexity of multi-task environments. This paper proposes a Real-Time Federated Multi-Task Learning framework using Small Language Models (RT-FMTL-SLM). We systematically evaluate five SLM architectures across benchmarks from the GLUE suite under multiple paradigm-complexity stages and data distributions. Our findings reveal that RT-FMTL-SLM achieves a per-client resource reduction of 74–93% relative to centralized training, though it incurs a 2.2–3.7× training time overhead due to real-time synchronization. We further analyze the trade-offs of Low-Rank Adaptation (LoRA), which reduces communication volume by over 90% but challenges regression stability in tasks like STS-B. DistilBERT emerges as the most robust architecture for multi-task stability under heterogeneous conditions. The source code is available at: https://github.com/vinh1988/RT-FMTL-SLM.git.**

***Index Terms*—Federated Learning, Small Language Models, Multi-Task Learning, Data Heterogeneity, Privacy-Preserving Machine Learning**

I. INTRODUCTION

With the proliferation of NLU systems on mobile devices, enterprise platforms, and edge services, the demand for privacy-preserving collaborative learning has grown at an alarming pace. In many practical scenarios, NLU data includes sensitive user information (e.g., personal messages, medical text, behavioral logs) making it challenging to collect centralized data due to privacy and regulatory constraints. FL resolves such shortcomings by storing raw data on client devices and only sharing model updates [1]. At the same time, distributed AI systems are expected to support many users, domains, and tasks in parallel [1], [2]. The trend makes centralized single-pipeline training less feasible and encourages federated training models, which can learn collaboratively without breaking the data locality.

While FL is a promising technique for NLU, there are three major concerns. First, client data usually is non-IID. Given that diverse users and organizations write language data with different vocabularies, label compositions and preferences, convergence is typically slowed, and global model quality is poor [9], [2]. Second, task heterogeneity is common in realistic deployments. A federation could be made with intent classification, sentiment analysis, question answering, or even generation-related goals all at once [2], [3], [5]. Classical one-model FL methods are not suited to this mixed-task regime. Third, communication and scalability are essential. In general, the repeated syncing of multiple clients can lead to significant overhead, particularly when the size of the model is large and the connected clients unstable [1], [3]. Recent literature on federated multi-task learning (e.g., MOCHA, FedBone, MIRA, encoder–decoder based M-Fed) indicates that data heterogeneity and task heterogeneity need to be considered as first-class design considerations [2], [3], [4], [5].

Small Language Models (SLMs) are especially appealing in federated systems because they improve efficiency without losing language pretraining functionality. Key advantages include:

- Communication efficiency: A smaller parameter set reduces the communication cost per round (both upload and download).
- Reduced computation cost: SLMs use fewer resources of local memory and training time, allowing for broader client involvement.
- Edge-device deployment: compact models are easier to deploy and update on resource-constrained devices.

This role translates into parameter-efficient adaptation techniques (e.g., adapter-style and low-rank updates) that preserve competitive performance while decreasing federated overhead.

The vast majority of existing FL pipelines are batch-based: clients train on multiple local steps and sync periodically [1], [9]. This design is strong on static datasets,

but it is not as well suited to streaming and dynamic NLU, in which patterns of language use, subjects, themes and users' intent evolve over time. Real-time federated learning seeks to reduce feedback loops by aiding in speeding up synchronization, ad-hoc adaptation, and more timely model updates. In online NLU applications (such as chatbots, live moderation, event-based analytics, etc.), latency of global updates can degrade performance very quickly. Thus, the implementation of real-time synchronization mechanisms seems to be on the rise for relevance and robustness of production systems.

Expanding on this motivation, and the reviewed literature, this work proposes:

- A real-time FL-MTL approach for heterogeneous NLU tasks.
- Small language models in federated settings to increase efficiency and deployability.
- Analysis of how data heterogeneity affects FL-MTL convergence and downstream task performance.
- Evaluation across different NLU benchmark tasks, local and global.

Herein, the rest of the paper is structured as follows. Section II reviews the relevant work in federated and multi-task learning. Section III defines the formal problem of RT-FMTL. In Section IV we detail the system architecture and optimization strategies. The experimental setup and results analysis are described in Section V. Last but not least, Section VI wraps up the paper and presents the next steps.

II. RELATED WORK

A. Federated Learning for Natural Language Processing

There are many examples of FL utilized in NLP/NLU applications such as text classification and personalized language services where raw text cannot be aggregated [1], [9]. Basic optimization methods like FedAvg had set up the typical client-update-server-aggregation loop, and FedProx added regularization to enable better behavior against client drift [1], [9]. Federated multi-task approaches, such as MOCHA, demonstrate that inter-task relatedness can be made more powerful under heterogeneous setups by improving statistical efficiency and robustness to system issues [2]. Newer approaches such as FedBone take this further by explicitly separating shared from task-dependent elements in large heterogeneous multi-task architectures [3]. Notwithstanding this progress, a large number of FL-for-NLP studies rely on limited task diversity or mainly rely solely on single-task user personalization.

B. Multi-task learning NLP

Multi-task learning (MTL) in NLP generally employs one of the two paradigms: hard parameter sharing and soft parameter sharing [2], [10]. Hard sharing employs a common encoder, instead of a separate task head for each instance of the model, saving the model from overfitting and training cost. Soft sharing retains task-specific factors but promotes similarity either by regularization or by limiting interaction. MTL improves representation learning by transfer of useful linguistic features between tasks. On the other hand, it also comes with known drawbacks, such as task imbalance (over-representation of dominant tasks with lower-level tasks) and negative transfer (one task harms others). These problems are more pronounced in the context of federated training because client involvement is partial, and the balance of data distribution between rounds is not consistent.

C. NLP and Small Language Models

SLMs like DistilBERT, MiniLM, and ALBERT have been shown that careful compression and distillation can preserve good performance with a cost between parameter count and inference [6], [7], [8]. The common reduction methods are knowledge distillation, parameter sharing, low-rank factorization, and architectural simplification. These capabilities make those SLMs very appropriate for edge and distributed learning: they decrease the transmission overhead, reduce on-device training cost, and achieve faster federated iteration cycles. In settings with multistask federated solutions, SLMs also facilitate scaling onto larger numbers of clients and tasks on fixed resource budgets for the system.

D. Federated Learning in Heterogeneous Data

Non-IID data is still one of the major bottlenecks in FL. Client-dependent language patterns, domain shifts, and label bias may lead to biased updates at the local level and may disrupt global optimization at the global level. This problem is aggravated in heterogeneous NLU deployments where the difference between client tasks exists (cross-client task diversity) whereby the clients can perform partial representations for optimizing different objectives. As previous work report, aggregation without attention to the structure of a task will degrade velocity and final generalization, whereas when heterogeneity is mixed up (thereby aggregation). Since convergence speed can harm fine generalization and final speed depending on task structure, techniques such as personalized layers, proximal regularization, gradient conflict reducing, and encoder–decoder decomposition have been proposed to bridge this gap [9], [10], [5].

E. Research Gap

Current studies offer powerful foundations, and still constitute, however, an important gap to be filled in terms of practical NLU systems:

- No real-time FL-based libraries: the majority of approaches only adopt offline versions for short time periods rather than for real-time adaptation.

- Inconsistent and partial inclusion of FL + MTL + SLM: the efficiency, heterogeneity and multi task-transfer of the models are frequently reported in different ways in papers.
- Insufficiently analysis heterogeneity in NLU FL: the present work tends to overlook how task diversity and the relationship between task diversity, data skew, and communication limitations work together.

This gap spurs a unified, real-time federated multi-task architecture comprising efficient language modeling and heterogeneous optimization for distributed NLU, coupled with optimization awareness.

III. PROBLEM DEFINITION

A. Formalization of Real-Time Federated Multi-Task Learning For NLU

We opt a real-time federated multi-task learning (RT-FMTL) environment based on natural language understanding (NLU) in which training goes over communication rounds for $t = 1, 2, \ldots$ [1], [2], [3]. Each round a subset of clients is selected, local updates are performed on private data streams, and the model gets pushed to a central server for aggregation. With regards to its distribution the client data distribution may vary over time, unlike in case of offline FL only, therefore, both local goals and optimum parameters are time bound. Let $\mathcal{C} = \{1, \ldots, C\}$ be the set of clients and $\mathcal{K} = \{1, \ldots, K\}$ be the set of NLU tasks. For each client $c \in \mathcal{C}$ it has only one major task $\tau(c) \in \mathcal{K}$ (or a small task subset), where new samples can continuously be added from a local stream. To learn a federated model that maintains cross-task transfer quality with the specific task quality in context of non-IID and time-varying data.

B. Data and Task Heterogeneity in Clients

Heterogeneity comes in two forms:

- Heterogeneity of data: the local distributions in clients c_i, c_j will be different (i.e. $P_{c_i}(x, y) \neq P_{c_j}(x, y)$ (label skew, domain shift, vocabulary differences).
- Task heterogeneity: clients may optimize different losses and output spaces according to $\tau(c)$, yielding dissimilar gradients across tasks.

Moreover, real-time environments have client availability and update frequency asynchronous, leading to system heterogeneity and stale updates [1], [9], [3]. Thus, the optimization process must account for partial participation and non-uniform local progress.

C. Optimization Objective for Federated Multi-Task Learning

Let $\ell_{\tau(c)}(\theta_e, \theta_{\tau(c)}; \xi)$ be the task-specific loss cost for sample $\xi \in \mathcal{D}_c^t$. The global objective at round t is:

$$\min_{\theta_e, \{\theta_k\}} F_t(\Theta) = \sum_{c=1}^{C} p_c^t \, \mathbb{E}_{\xi \sim \mathcal{D}_c^t} \big[\ell_{\tau(c)}(\theta_e, \theta_{\tau(c)}; \xi)\big] + \lambda \, \Omega(\{\theta_k\}), \tag{1}$$

where p_c^t represents the weight of the aggregation (e.g., the weight per unit local sample), Ω is a regularizer to promote a stable cross-task learning and $\lambda \geq 0$ controls the sharing-specialization trade-off. After localized updates, the server aggregates common encoder parameters by weighted averaging,

$$\theta_e^{t+1} = \sum_{c \in \mathcal{C}_t} \alpha_c^t \, \theta_{e,c}^{t+1}. \tag{2}$$

and updates custom task-specific parameters, as a result of task-wise aggregation,

$$\theta_k^{t+1} = \sum_{c \in \mathcal{C}_t \colon \tau(c) = k} \beta_c^t \, \theta_{k,c}^{t+1}. \tag{3}$$

In $\mathcal{C}_t$, the client subset was measured at round t and α_c^t, β_c^t are normalized aggregation coefficients [1], [2], [5]. With privacy and heterogeneity requirements in place, the goal captures RT-FMTL in real time by jointly optimizing both shared representations and task-specific parameters.

IV. PROPOSED METHOD

A. System Overview

In this work, we propose a multistructure federated multi-task framework designed for real-time NLU applications. As illustrated in Fig. 1, a high level orchestration server, along with multiple, heterogeneous client nodes, structure the system. A parameter server is responsible for maintaining the global state at the center and controls the Model Aggregation logic. The communication relies on low-latency WebSockets for a bidirectional exchange of model updates and aggregated parameters.

Federated components and privacy constraints of a system are based on:

- Global server $\mathcal{S}$: coordinates client sampling, model broadcast, and aggregation.
- Distributed clients $\mathcal{C}$: edge devices/silos to store the raw text.
- Task set $\mathcal{K}$: the heterogeneous NLU tasks (intent classification, sentiment analysis, sequence labeling).
- Local datasets $\mathcal{D}_c^t$: single private, time-stamped data for client c at round t that is private.

Samples are never uploaded, for privacy purposes. Each client transmits only model updates (or compressed

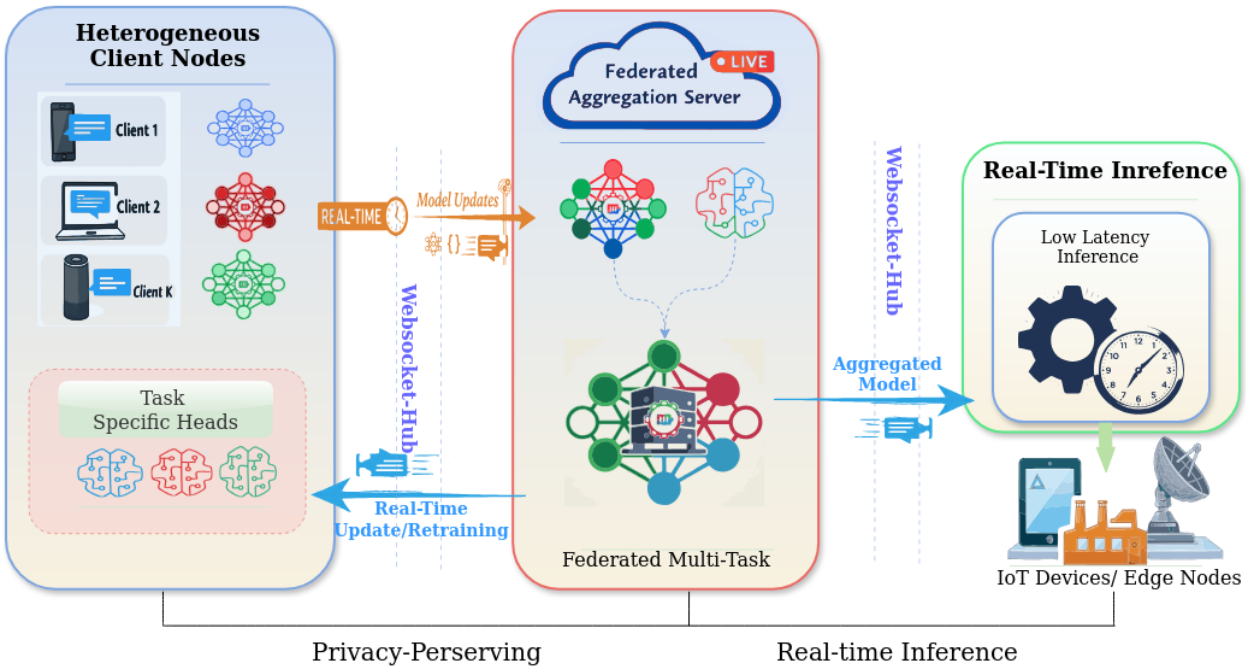

Fig. 1: Architecture of the Proposed Real-time Federated Multi-Task Learning Framework. The system enables seamless integration of streaming data processing and collaborative model refinement.

adapter and head updates) and serves as host-side aggregation [1], [4]. This is consistent with the usual federated privacy model in which data locality is preserved and collaborative training is promoted in the field.

B. Backbone of the Small Language Model

This framework is adapted with specialized Small Language Models (SLMs) to trade-off computational efficiency and language fluency. Our current solution also supports DistilBERT [6], BERT-Medium [14], MiniLM [7] architectures. These are shared feature extractors (encoder) that are fine-tuned collaboratively across the federated network. Our SLMs drastically minimize the communication overhead in contrast to standard-sized transformers.

C. Multi-Task Learning Strategies

We use a hybrid strategy of parameter-sharing. The lower layers of the transformer (encoder) are the same across the tasks that represent linguistic representation universally. In contrast, the top layers are task-specific classification or regression heads. We thus do an adaptive loss weighting mechanism to reduce the negative transfer:

$$\mathcal{L}_{total} = \sum_{j=1}^{T} \alpha_j \mathcal{L}_j \tag{4}$$

where α_j is dynamically tuned according to the convergence rate of each task j.

We define the parameters

$$\Theta = \{\theta_e, \{\theta_k\}_{k=1}^{K}\}, \tag{5}$$

for which θ_e is a common small-language-model encoder and θ_k is the task-based prediction head of task k [3], [5], [4]. Prediction is for an input-output pair (x, y) given task k.

$$\hat{y} = h_k\big(f_e(x;\theta_e);\theta_k\big), \tag{6}$$

where $f_e(\cdot)$ represents the shared encoder and $h_k(\cdot)$ refers to the task head. It allows for parameter sharing amongst transferable linguistic traits but maintains task-level specialization.

D. Federated Optimization and Real-time adaptation

The process of federated optimization is detailed in Algorithm 1, which describes the federated optimization of the system. Clients will be able to update in real-time while the clients are still synchronizing updates asynchronously through WebSockets using this system to adapt in real-time. Updating the server aggregates updates with an example of a task-aware version of FedAvg (i.e. a version of FedAvg) and rounds off the updating with noise from heterogeneous data sources (Non-IID), but regularizes noise correctly.

E. Communication-Efficient Fine-Tuning (LoRA)

To optimize resource utilization on edge devices even further, we implement Low-Rank Adaptation (LoRA)

Algorithm 1 RT-FMTL: Real-Time Federated Multi-Task Learning

Server Initialize: Shared core ϕ_0, Task heads $\{\psi_{j,0}\}_{j=1}^{T}$
for each round $t = 1, 2, \ldots$ **do**
  Server transmits $\phi_t, \psi_{Task(k),t}$ to client $k \in K$ via WebSockets
  for each client $k \in K$ **in parallel do**
    Local training: $\{\Delta\phi_k, \Delta\psi_{Task(k),k}\} \leftarrow \text{LocalUpdate}(\mathcal{D}_k, \phi_t, \psi_{Task(k),t})$
    Updates $\{\Delta\phi_k, \Delta\psi_{Task(k),k}\}$ sent to Server using WebSockets
  end for
  Task-Aware Aggregation:
    $\phi_{t+1} \leftarrow \phi_t + \eta \sum_{k=1}^{K} \frac{n_k}{N} \Delta\phi_k$ {Universal Knowledge}
  for each task $j \in \{1, \ldots, T\}$ **in parallel do**
    $\psi_{j,t+1} \leftarrow \psi_{j,t} + \eta \sum_{k \in K_j} \frac{n_k}{N_j} \Delta\psi_{j,k}$ {Heads Specialization}
  end for
  Server redistributes $\{\phi_{t+1}, \psi_{Task(k),t+1}\}$ for Local Validation
end for

[12]. This includes when to update in local B, where instead of updating the full weight matrices $W \in \mathbb{R}^{d\times d}$, we optimize the low rank decomposition matrices A and B, where:

$$W' = W + \Delta W = W + BA. \tag{7}$$

$A \in \mathbb{R}^{r\times d}$ and $B \in \mathbb{R}^{d\times r}$ with rank $r \ll d$. For big updates, only the matrices A and B are sent, minimizing communication volume to up to 99%.

V. EXPERIMENTS

A. Experimental Setup

1) Datasets and Tasks: The evaluation employs three core NLU benchmarks representing diverse natural language understanding challenges from the GLUE benchmark [11].

SST-2 [15] is a binary classification task for sentiment polarity detection, comprising 66,477 training, 872 validation, and 1,821 test samples. Performance is measured by accuracy and F1-score [20].

QQP [16] is a binary classification task for paraphrase detection, with 323,415 training, 40,431 validation, and 39,081 test samples. Evaluation uses accuracy and F1-score [20].

STS-B [17] is a regression task for semantic similarity scoring over sentence pairs on a 1–5 scale, with 4,249 training, 1,500 validation, and 1,379 test samples. Performance is measured by Pearson [18] and Spearman correlation [19].

TABLE I: Experimental Design Matrix

Exp.	Paradigm	Task Config.	Clients	Data Dist.
1	Centralized MTL	Multi-Task	3	Full access
2	Centralized Single	SST-2 only	1	Full access
3	Centralized Single	QQP only	1	Full access
4	Centralized Single	STS-B only	1	Full access
5	Federated MTL	Multi-Task	3	IID
6	Federated MTL	Multi-Task	3	IID + LoRA
7	Federated MTL	Multi-Task	9	Non-IID (3×3)
8	Federated Single	QQP only	3	Non-IID
9	Federated Single	SST-2 only	3	Non-IID
10	Federated Single	STS-B only	3	Non-IID

2) Federated Data Partitioning: IID Partitioning distributes data uniformly across clients (Experiments 5-6 in Table I, 3 clients, multi-task scenarios).

Non-IID Partitioning uses a Dirichlet distribution (α=0.5) to simulate realistic data heterogeneity, with 3 clients per task (9 total), yielding 107,805 / 22,159 / 1,416 samples per client for QQP, SST-2, and STS-B respectively (Experiments 7-10 in Table I).

Validation sets are *not* split across clients in any paradigm, ensuring consistent global model assessment.

3) Baselines: Three baselines are evaluated: (1) Centralized MTL: all 5 SLM variants trained on centralized data with shared encoder and task-specific heads, serving as the upper-bound reference (Exp. 1 in Table I); (2) Federated Single-Task: per-task federated experiments with 3 clients and non-IID data (Exps. 8-10 in Table I); and (3) Large-model federated approaches: comparison against larger LMs to assess communication efficiency and performance trade-offs.

4) Evaluation Metrics: Task-level metrics are accuracy & F1-score [20] (SST-2, QQP) and Pearson [18] / Spearman correlation [19] (STS-B) [11]. System-level federated metrics include communication cost per round, convergence speed, and total training time.

5) Implementation Details: All experiments use learning rate 2×10^{-5} with 1,000-step warmup, batch size 8 (centralized) or 8-16 (federated), weight decay 0.01, gradient clipping (max-norm 1.0), and random seed 42. Centralized runs train for 25 epochs; federated clients perform 1 local epoch per round. Communication uses ports 8771-8775 with 60 s client timeout and 3,400 s round timeout, plus a 3-retry exponential-backoff mechanism.

B. Results and Discussion

This section evaluates the proposed Federated MTL framework across five thematic dimensions: overall baseline comparisons, the impact of MTL, paradigm performance gaps, real-time suitability, and communication efficiency.

TABLE II: Comparative Analysis of SLM Performance across Paradigms

Model	Paradigm	Task	Dist.	Tuning	SST-2		QQP		STS-B		System	System
					Acc.	F1	Acc.	F1	Pear.	Spea.	Time (s)	Res (GB)
DistilBERT [6]	FL	ST	Non-IID	FFT	—	—	0.8408	0.7724	—	—	34904.44	1.0956
	FL	ST	Non-IID	FFT	—	—	—	—	0.7753	0.7421	15452.52	1.0988
	FL	ST	Non-IID	FFT	0.9346	0.9392	—	—	—	—	23586.44	1.0972
	FL	**MTL (ours)**	IID	LoRA	0.8257	0.8433	0.7535	0.6754	0.5108	0.4919	21245.26	0.2852
	FL	**MTL (ours)**	Non-IID	FFT	0.9232	0.9298	0.8579	0.8054	0.6963	0.6967	52903.97	1.0968
	FL	**MTL (ours)**	IID	FFT	**0.9461**	**0.9514**	0.8913	0.8610	0.8284	0.8136	56454.46	1.0896
	Cent.	ST	IID	FFT	—	—	—	—	**0.8712**	**0.8691**	4561.78	4.2470
	Cent.	ST	IID	FFT	0.9002	0.9061	—	—	—	—	18514.82	4.2460
	Cent.	ST	IID	FFT	—	—	0.8456	**0.8934**	—	—	4290.59	4.2547
	Cent.	MTL	IID	FFT	0.9037	0.9077	0.8137	0.8778	0.8635	0.8614	24449.68	4.3214
BERT-Medium [14]	FL	ST	Non-IID	FFT	—	—	0.7989	0.6852	—	—	23880.89	0.6789
	FL	ST	Non-IID	FFT	—	—	—	—	0.6521	0.6212	10261.92	0.6810
	FL	ST	Non-IID	FFT	0.9300	0.9348	—	—	—	—	15439.04	0.6786
	FL	**MTL (ours)**	Non-IID	FFT	0.8933	0.8988	0.8183	0.7441	0.4219	0.4132	33374.50	0.6790
	FL	**MTL (ours)**	IID	FFT	0.9392	0.9450	0.8962	0.8655	0.8258	0.8052	45258.21	0.6807
	FL	**MTL (ours)**	IID	LoRA	0.8131	0.8362	0.7388	0.6805	0.6823	0.6512	12885.74	0.1840
	Cent.	ST	IID	FFT	—	—	—	—	0.8673	0.8648	3381.83	2.6467
	Cent.	ST	IID	FFT	0.8933	0.8970	—	—	—	—	13492.62	2.6711
	Cent.	ST	IID	FFT	—	—	0.8358	0.8859	—	—	3205.51	2.6467
	Cent.	MTL	IID	FFT	0.8865	0.8884	0.8333	0.8777	0.8615	0.8584	19257.40	2.6671
BERT-Mini [14]	FL	ST	Non-IID	FFT	—	—	0.7674	0.6201	—	—	9642.59	0.1969
	FL	ST	Non-IID	FFT	—	—	—	—	0.6504	0.6150	2547.10	0.1969
	FL	ST	Non-IID	FFT	0.9151	0.9196	—	—	—	—	4552.75	0.1969
	FL	**MTL (ours)**	Non-IID	FFT	0.8234	0.8226	0.7748	0.6668	0.4009	0.3939	10317.55	0.1969
	FL	**MTL (ours)**	IID	LoRA	0.7649	0.7915	0.6958	0.5602	0.2824	0.2773	4191.67	0.0630
	FL	**MTL (ours)**	IID	FFT	0.9266	0.9329	0.8350	0.7981	0.6940	0.6550	13341.01	0.1969
	Cent.	ST	IID	FFT	—	—	—	—	**0.8316**	**0.8283**	411.45	2.4487
	Cent.	ST	IID	FFT	0.8727	0.8773	—	—	—	—	3348.26	2.4487
	Cent.	ST	IID	FFT	—	—	0.7990	0.8536	—	—	318.61	2.4487
	Cent.	MTL	IID	FFT	0.8589	0.8650	0.7770	0.8580	0.7989	0.7973	3652.50	2.4651
MiniLM [7]	FL	ST	Non-IID	FFT	—	—	0.8421	0.7644	—	—	14583.05	0.3842
	FL	ST	Non-IID	FFT	—	—	—	—	0.7877	0.7489	4920.48	0.3840
	FL	ST	Non-IID	FFT	0.9243	0.9293	—	—	—	—	8447.20	0.3836
	FL	**MTL (ours)**	IID	FFT	0.9369	0.9428	**0.8966**	0.8662	0.8384	0.8197	27930.06	0.3835
	FL	**MTL (ours)**	Non-IID	FFT	0.8888	0.8974	0.8629	0.8150	0.7003	0.7061	20765.57	0.3846
	FL	**MTL (ours)**	IID	LoRA	0.7787	0.8102	0.7271	0.6060	0.0703	0.0621	18808.08	0.1090
	Cent.	ST	IID	FFT	—	—	—	—	0.8620	0.8592	2696.27	1.4929
	Cent.	ST	IID	FFT	0.9014	0.9051	—	—	—	—	8334.10	1.4950
	Cent.	ST	IID	FFT	—	—	0.7990	0.8656	—	—	2388.13	1.5067
	Cent.	MTL	IID	FFT	0.8796	0.8829	0.7917	0.8649	0.8580	0.8560	12437.12	1.5045
TinyBERT [13]	FL	ST	Non-IID	FFT	—	—	0.7010	0.3887	—	—	6226.63	0.0881
	FL	ST	Non-IID	FFT	—	—	—	—	0.2186	0.1755	1182.03	0.0881
	FL	ST	Non-IID	FFT	0.9060	0.9122	—	—	—	—	2514.06	0.0881
	FL	**MTL (ours)**	IID	FFT	0.9014	0.9145	0.7950	0.7494	0.5212	0.4754	6778.00	0.0881
	FL	**MTL (ours)**	Non-IID	FFT	0.7810	0.7782	0.6895	0.4733	0.0081	0.0063	6975.28	0.0881
	FL	**MTL (ours)**	IID	LoRA	0.7569	0.7902	0.6766	0.4475	0.3636	0.3505	3560.68	0.0356
	Cent.	ST	IID	FFT	—	—	—	—	0.7270	0.7230	2337.42	0.3427
	Cent.	ST	IID	FFT	0.8314	0.8358	—	—	—	—	4857.08	0.3427
	Cent.	ST	IID	FFT	—	—	0.7157	0.8139	—	—	2315.39	0.3427
	Cent.	MTL	IID	FFT	0.8211	0.8263	0.7230	0.8101	0.6474	0.6426	9227.55	0.3440

Note: Bold values indicate highest performance. For DistilBERT, red color indicates best proposed method (IID), dark yellow color indicates second best (Non-IID). ST = Single-Task, MTL = Multi-Task Learning, Cent. = Centralized, LoRA = Low-Rank Adaptation, FFT = Full Fine-Tuning, Pear. = Pearson Correlation, Spea. = Spearman Correlation.

1) Overall Performance (Comparison across Baselines): The comprehensive results in Table II demonstrate that Small Language Models (SLMs) integrated within a federated multi-task framework maintain high NLU performance across diverse tasks. Centralized MTL serves as the performance upper bound, yet DistilBERT under Federated MTL (IID) surprisingly achieves 0.9461 accuracy on SST-2, slightly outperforming its centralized counterpart. MiniLM also exhibits robust performance, reaching 0.8966 accuracy on QQP in the Federated MTL setting. These results validate that despite the inherent constraints of FL, SLMs can achieve competitive results relative to idealized centralized training.

2) Effect of Multi-Task Learning: MTL provides significant advantages in federated settings by leveraging shared representations across related tasks. As shown in Figure 2 and Table II, comparing FL-Single and FL-MTL (Non-IID) reveals complex task interactions. While SST-2 suffers modest negative transfer in certain configurations, regression tasks like STS-B often suffer more pronounced drops under Non-IID conditions. However,

classification tasks generally demonstrate better positive transfer; for instance, QQP benefits from the features learned via SST-2. Larger SLMs like DistilBERT and BERT-Medium show higher resilience to negative transfer due to their greater representational capacity.

3) Federated Learning Performance: The transition from centralized to federated learning introduces a "Reliability Gap," characterized by a moderate degradation ($\leq 5\%$) in classification accuracy for tasks like SST-2 and QQP. This gap is most visible in the STS-B Pearson correlation, which drops significantly for the smallest architectures like TinyBERT under Non-IID MTL conditions. Figure 2 maps this trajectory, showing how the proposed Federated MTL framework (highlighted in red for IID and dark yellow for Non-IID for the DistilBERT model in Table II) manages to close this gap by stabilizing performance across tasks.

Quantitatively, the red metrics (IID) represent the framework's optimal performance, showing that MTL can yield better results than isolated centralized training for specific models (e.g., DistilBERT reaching 0.9461 on SST-2). The dark yellow metrics (Non-IID) highlight the framework's resilience under data heterogeneity; while classification tasks like SST-2 maintain stable performance (e.g., 0.9232 for DistilBERT), regression tasks like STS-B exhibit greater sensitivity, dropping from 0.8258 to 0.4219 in BERT-Medium. This demonstrates that while the framework is remarkably robust for categorical benchmarks, regression stability is heavily dependent on model capacity and data distribution.

4) Real-Time Inference Analysis: For deployment on edge devices where real-time inference is critical, resource consumption and model footprint are the defining metrics. Figure 3 and Table III highlight that federated learning reduces individual client GPU memory demands by 77-93% relative to centralized training. TinyBERT and BERT-Mini emerge as the leading candidates for extremely resource-constrained devices (< 0.1 GB and < 0.2 GB per client respectively), albeit with bounded performance on complex tasks. MiniLM and DistilBERT provide a superior performance-efficiency balance for most real-time scenarios, maintaining high accuracy while keeping per-client resource usage significantly lower than the centralized baseline.

5) Communication Efficiency: Communication overhead is the primary bottleneck in federated MTL, with wall-clock training times increasing by 2.2-3.7$\times$ compared to centralized training (Table IV). However, as illustrated in Figure 4, the integration of parameter-efficient fine-tuning via LoRA [12] substantially mitigates this cost. LoRA reduces federated training time by 38-62% for most models by transmitting only low-rank adapter weights instead of full model parameters. This optimization enables the framework to maintain competitive accuracy while drastically reducing the bandwidth requirements, making it suitable for bandwidth-limited edge deployments.

VI. CONCLUSION

In this work, we presented a comprehensive investigation into real-time federated multi-task learning using Small Language Models across diverse Natural Language Understanding benchmarks. Our systematic evaluation across five architectures provided critical insights for the deployment of decentralized NLU systems. Our results demonstrate that federated learning effectively preserves data privacy while significantly optimizing resource distribution, achieving up to 93% reduction in per-client GPU resource consumption. However, the transition to real-time federated multi-task environments introduces notable synchronization overheads, as observed in our study using WebSocket-based communication.

A key challenge identified in our work is the impact of task heterogeneity. We found substantial negative transfer effects in regression tasks (STS-B) under Non-IID distributions, particularly for the smallest model variants (TinyBERT and BERT-Mini). Heterogeneity-aware aggregation and Low-Rank Adaptation (LoRA) proved essential for communication efficiency, though LoRA involves a trade-off in performance stability for complex tasks. Among the investigated SLM architectures, DistilBERT consistently offered the most robust performance, showing superior resilience to the negative transfer effects inherent in multi-task federated learning. For scenarios where memory is the primary bottleneck, architectures like MiniLM provide a viable alternative, though with less stability in regression regimes.

Future research will focus on developing adaptive task-sharing strategies to further mitigate negative transfer in multi-head architectures. Additionally, we plan to explore personalized federated learning techniques and the integration of differentially private mechanisms within the real-time communication layer, ensuring that efficiency gains do not compromise privacy in production NLU deployments.

REFERENCES

[1] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas. Communication-efficient learning of deep networks from decentralized data. In *AISTATS*, 2017.

[2] V. Smith, C.-K. Chiang, M. Sanjabi, and A. Talwalkar. Federated multi-task learning. In *NeurIPS*, 2017.

[3] Y. Chen, T. Zhang, X. Jiang, Q. Chen, C. Gao, and W. Huang. FedBone: Towards large-scale federated multi-task learning. *arXiv preprint arXiv:2306.17465*, 2024.

[4] A. Elbakary, C. Ben Issaid, T. ElBatt, K. Seddik, and M. Bennis. MIRA: A method of federated multi-task learning for large language models. *arXiv preprint*, 2025.

[5] J. Zhou, W. Bao, J. Wang, D. Zhang, X. Zhang, and Y. Zhang. Multi-task federated learning with encoder–decoder structure. *arXiv preprint arXiv:2504.09800*, 2025.

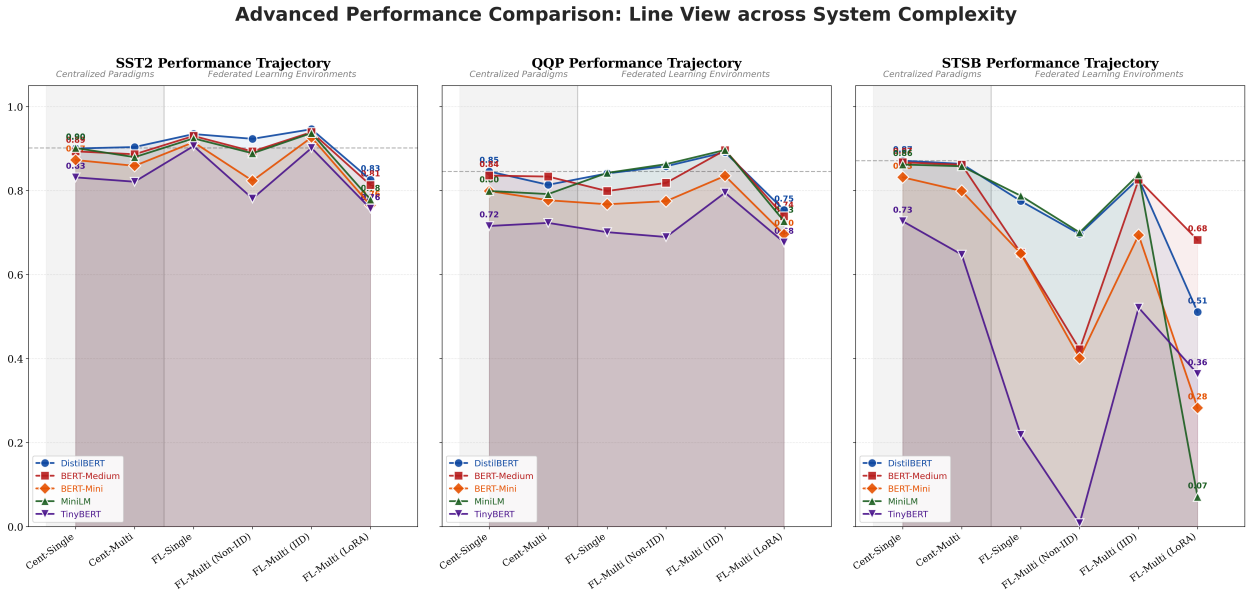

Fig. 2: Performance Trajectory across six paradigm-complexity stages for all five SLM architectures. Dashed vertical lines separate centralized from federated settings. The final point (LoRA) illustrates performance recovery under communication-efficient fine-tuning despite maximal system complexity.

TABLE III: Resource Usage Trajectory (GB, per-client average)

Model	Cent-S	Cent-M	FL-S	FL-M (Non-IID)	FL-M (IID)	FL-M (LoRA)
DistilBERT	4.2492	4.3214	1.0972	1.0968	1.0896	0.2852
BERT-Medium	2.6549	2.6671	0.6795	0.6790	0.6807	0.1840
BERT-Mini	2.4487	2.4651	0.1969	0.1969	0.1969	0.0630
MiniLM	1.4982	1.5045	0.3839	0.3846	0.3835	0.1090
TinyBERT	0.3427	0.3440	0.0881	0.0881	0.0881	0.0356

TABLE IV: Training Time Trajectory (seconds)

Model	Cent-S	Cent-M	FL-S	FL-M (Non-IID)	FL-M (IID)	FL-M (LoRA)
DistilBERT	9,122	24,450	24,648	52,904	56,454	21,245
BERT-Medium	6,693	19,257	16,527	33,375	45,258	12,886
BERT-Mini	1,359	3,653	5,581	10,318	13,341	4,192
MiniLM	4,473	12,437	9,317	20,766	27,930	18,808
TinyBERT	3,170	9,228	3,308	6,975	6,778	3,561

[6] V. Sanh, L. Debut, J. Chaumond, and T. Wolf. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.

[7] W. Wang, F. Wei, L. Dong, H. Bao, N. Yang, and M. Zhou. MiniLM: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. In *NeurIPS*, 2020.

[8] Z. Lan, M. Chen, S. Goodman, K. Gimpel, P. Sharma, and R. Soricut. ALBERT: A Lite BERT for self-supervised learning of language representations. In *ICLR*, 2020.

[9] T. Li, A. K. Sahu, M. Zaheer, M. Sanjabi, A. Talwalkar, and V. Smith. Federated optimization in heterogeneous networks. In *MLSys*, 2020.

[10] T. Yu, S. Kumar, A. Gupta, S. Levine, K. Hausman, and C. Finn. Gradient surgery for multi-task learning. *arXiv preprint arXiv:2001.06782*, 2020.

[11] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman. GLUE: A multi-task benchmark and analysis platform for natural language understanding. In *ICLR*, 2019.

[12] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen. LoRA: Low-rank adaptation of large language models. In *ICLR*, 2022.

[13] X. Jiao, Y. Yin, L. Shang, X. Jiang, X. Chen, L. Li, F. Wang, and Q. Liu. TinyBERT: Distilling BERT for natural language understanding. In *Findings of EMNLP*, 2020.

[14] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *NAACL*, 2019.

[15] R. Socher, A. Perelygin, J. Wu, J. Chuang, C. D. Manning, A. Ng, and C. Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *EMNLP*, 2013.

[16] L. Sharma, L. Graesser, N. Nangia, and U. Evci. Natural language understanding with the quora question pairs dataset. *arXiv preprint arXiv:1907.01041*, 2019.

[17] D. Cer, M. Diab, E. Agirre, I. Lopez-Gazpio, and L. Specia. SemEval-2017 Task 1: Semantic textual similarity multilingual and crosslingual focused evaluation. In *SemEval*, 2017.

[18] K. Pearson. VII. Note on regression and inheritance in the case of two parents. *Proceedings of the Royal Society of London*, 58:240–

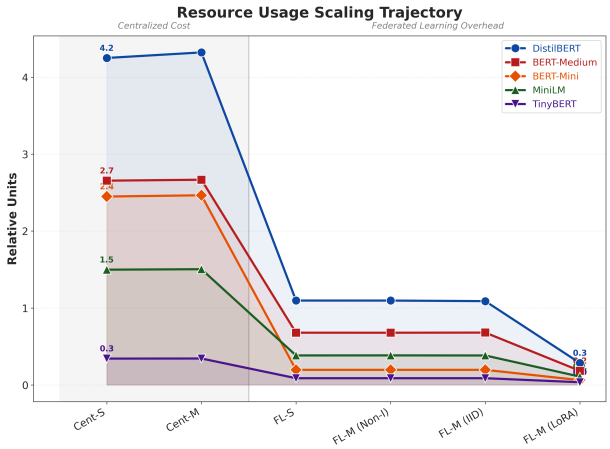

Fig. 3: Resource Usage Trajectory. Federated settings produce a clear 3-4× reduction in per-client resource usage relative to centralized training because each client holds only a data partition. LoRA further halves resource demands.

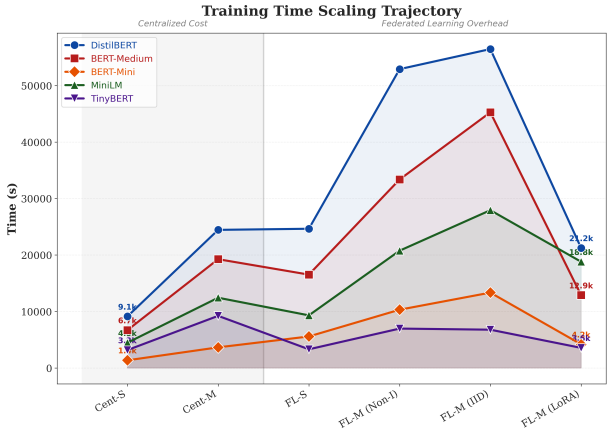

Fig. 4: Training Time Trajectory capturing the wall-clock training durations. The transition from centralized to federated multi-task settings introduces a significant overhead, which LoRA substantially offsets by limiting parameter transmission.

242, 1895.

[19] C. Spearman. The proof and measurement of association between two things. *The American Journal of Psychology*, 15(1):72–101, 1904.

[20] C. J. Van Rijsbergen. *Information Retrieval*. Butterworth-Heinemann, 2nd edition, 1979.